

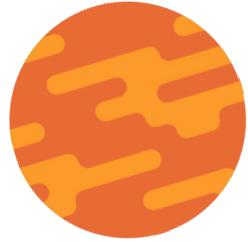
DataStax

Developers

Hands-On: Real-time
graph ETL for modern data pipelines with

Quine & Apache Cassandra®





David
Dieruf

Rags
Srinivas

Artem
Chebotko

Stefano
Lottini

Aleksandr
Volochnev

Aaron
Ploetz



Jack
Fryer

Kirsten
Hunter

Cedrick
Lunven

Mary
Grygleski

Ryan
Welford

David
Gilardi



DataStax Developers Crew

Developer Advocate



DataStax



Aaron Ploetz



@aaronploetz



@aploetz



@aploetz

- Former SWE/DevOps/DB Lead @ **GRAINGER** & **TARGET**.
- Host - Apache Cassandra Corner podcast
- Worked as an author on:
 - Mastering Apache Cassandra 3.x
 - Seven NoSQL Databases in a Week



Developer Relations



that Dot



@michael-aglietti



@maglietti

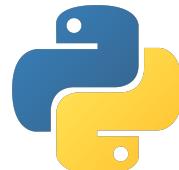


@michaelaglietti

- DevRel Leader
- Data-driven fanatic
- Solutions Architect
- Systems Engineer
- Information Architecture
- LAMP stack developer

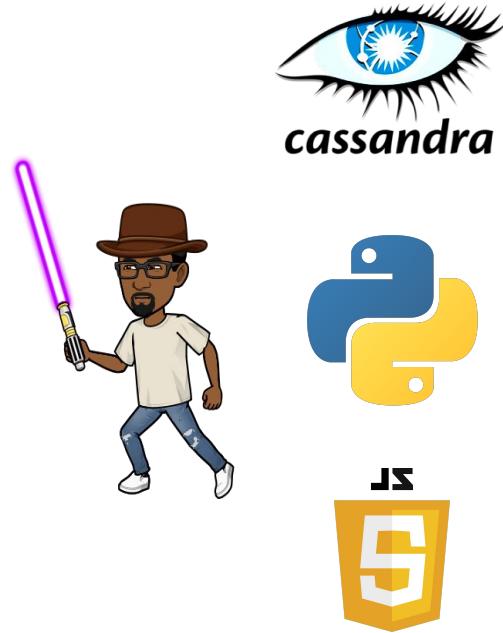
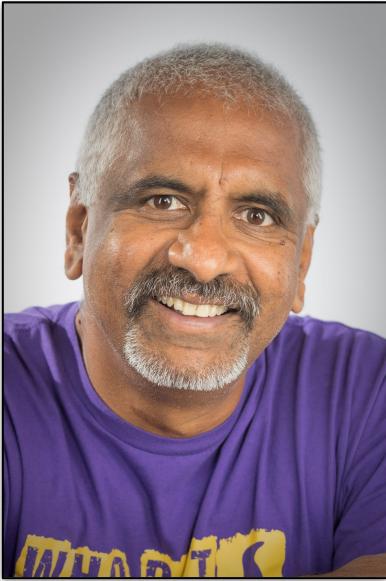


Quine



Michael Aglietti

Developer Advocate



- Developer/Architect
- Mechanical Engineer (so many moons ago)
- Distributed systems
- Love to teach and communicate
- Inner loop == developer productivity



Raghavan "Rags" Srinivas



@rags



@ragsns



@ragss



And a badge for your pleasure!

01



Introduction

02



Quine streaming graph



Use Cases

04



cassandra

Architecture

05



Hands on

06

What's next? Quiz, Homework, Next week

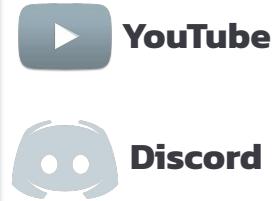
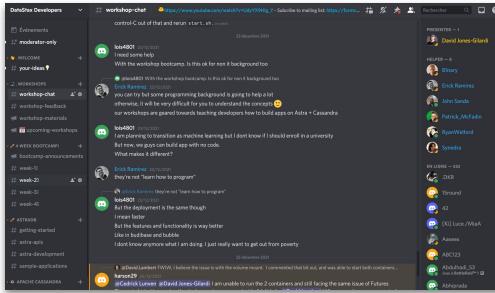


Agenda

Livestream: youtube.com/DataStaxDevs

Questions: <https://dtsx.io/discord>

Agenda



Games and quizzes: menti.com

How much experience do you have with the Spring Framework ?



1

Attend the live sessions

Database + GraphQL + PlayGround



DataStax
Astra DB



The screenshot shows the Gitpod IDE interface. On the left is the Explorer sidebar with project files like 'WORKSHOP-SPRING-STARGATE' and 'ManagesDemoApplication.java'. The main area shows code snippets for 'StargateDemoApplication.java' and 'ManagesDemoApplication.java'. Below the code are tabs for 'PROBLEMS', 'OUTPUT', 'TERMINAL', and 'DEBUG CONSOLE'. The terminal shows 'gitpod workspace/workspace/spring-stargate \$'. To the right are logos for 'npm', 'node.js', 'Maven', and 'Gitpod'.

The screenshot shows the Quine interface. It features a 'Real-time Graph ETL' section with a graph visualization showing nodes and edges. Below it is a 'Graph' section with a detailed view of a specific node's connections and timestamped events. The top navigation bar includes 'Docs', 'Getting Started', and 'Recipes'.

Quine



The screenshot shows a GitHub repository page for 'DataStax-Examples / todo-astra-jamstack-netlify'. It displays the repository structure, commit history, and release information. The GitHub logo is centered below the repository details.

GitHub

2

Complete Workshops Labs

Quine



cassandra



Today's Workshop: Quine & Cassandra

01



Introduction

02



Quine streaming
graph

03



Use Cases

04



Architecture

05



Hands on

06

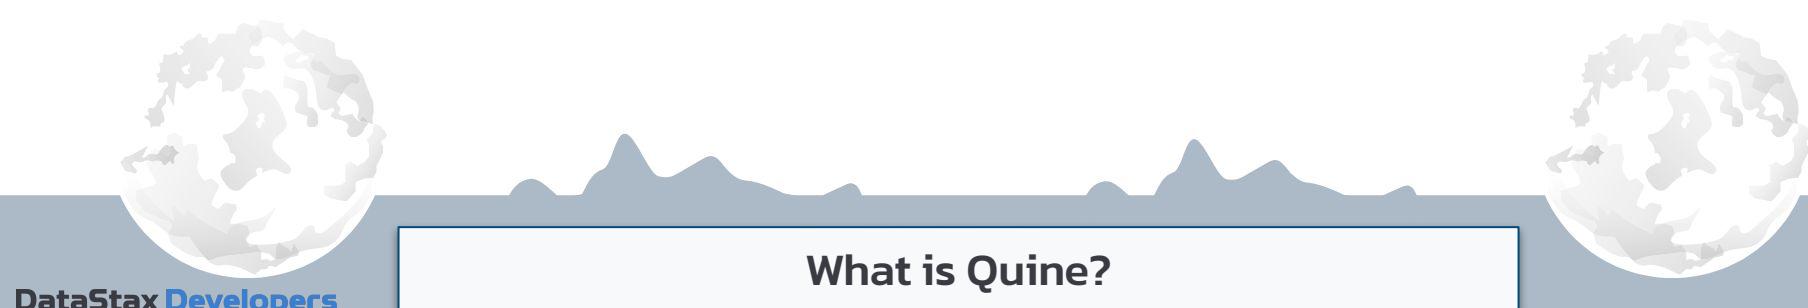
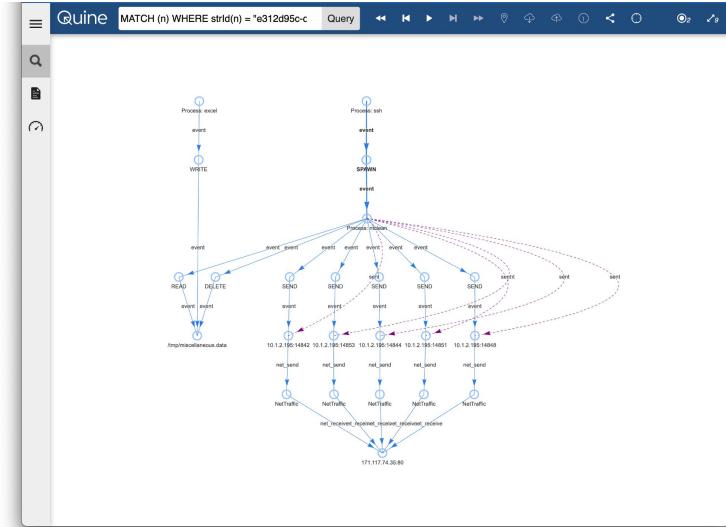
What's next?
Quiz, Homework, Next week



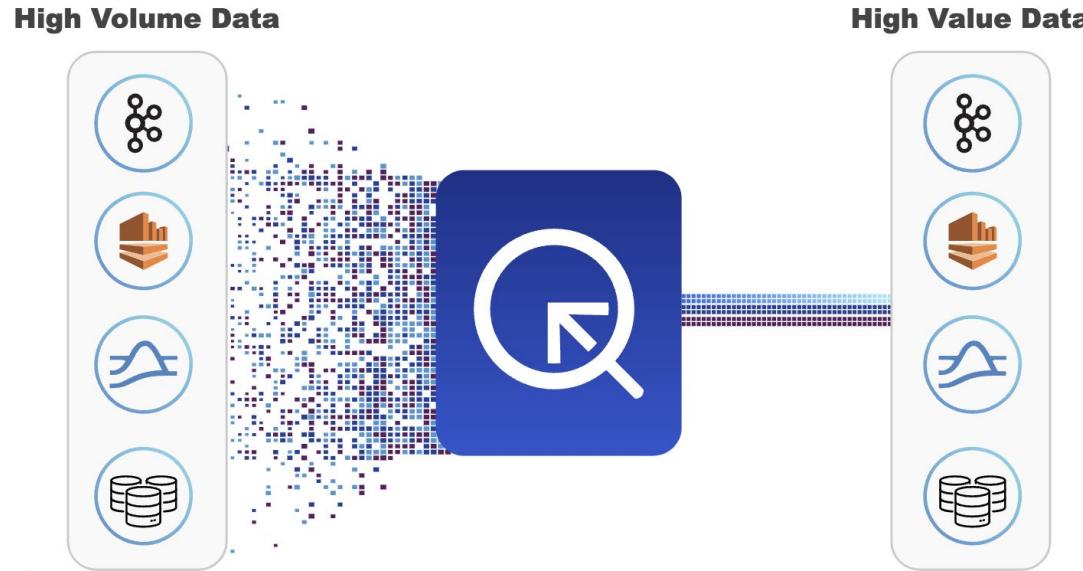
Agenda

Quine

- Open-source Graph ETL platform
- Graph data model easily joins multiple data sets
 - Track complex categorical relationships
 - No time windows to manage
 - Automated out-of-order data handling
 - High cardinality data sets
 - Flexible data structure
- Standing queries operate on data in real-time, or to generate events



Architecturally, Quine fits into a streaming event pipeline in between streaming event processors



Quine can shape, filter, analyze, or take action based on patterns matched in the event stream

High-Level Architecture

01



Introduction

02



Quine streaming
graph

03



Use Cases

04



Architecture

05



Hands on

06

What's next?
Quiz, Homework, Next week



Agenda



Log Analysis

General IT

Combine logs from multiple services to understand system performance issues and their root causes.

Streaming ETL

Event Streaming Architecture

Real-time ETL on streaming data, when data is received from multiple sources, requires new tools to accommodate out-of-order data.

Video Observability

Online Video

Video delivery monitoring is hard. Logs come from multiple sources with high dimensionality and cardinality, at scale.

Fraud Detection

Finance

Financial fraud detection typically relies on analysis of event volumes and transactions bypassing much of the categorical data.

Auth Fraud

General IT

Low frequency authentication attacks, such as password spraying and credential stuffing, circumvent authentication policies.

Blockchain Fraud

Finance - Blockchain

Monitor, detect and trace fraudulent transactions across blockchains in real time.



Quine

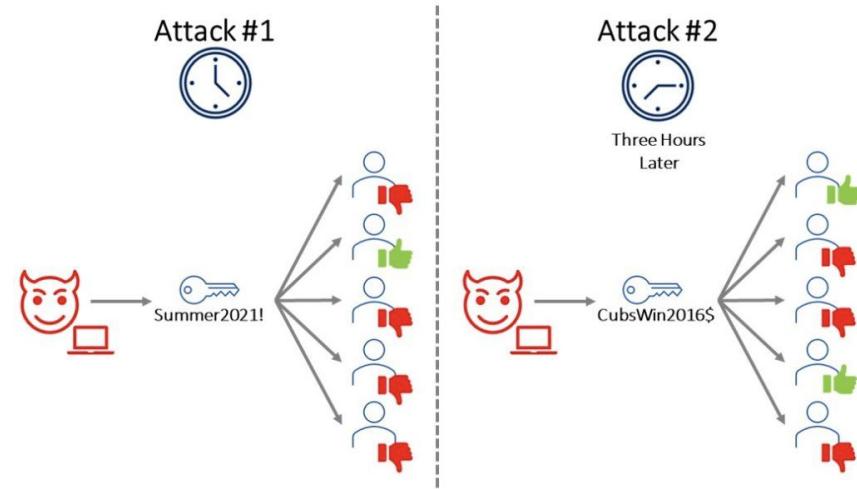


Value for Data Driven Businesses

Password Spraying

According to [Microsoft](#), the three steps to conduct a password-spraying attack are:

- **Acquire a list of usernames:** starting with a list of names: `firstname.lastname@company.com`
- **Spray passwords:** testing popular and common passwords (e.g., abc123, password, and Summer21!). See the [top 10,000 passwords](#).)
- **Gain access:** one of the tested attempts works, and the account can be abused to enumerate assets in the AD network, exploit authenticated services and put the organization at risk.



What is Password Spraying?

01



Introduction

02



Quine streaming
graph

03



Use Cases

04



Architecture

05



Hands on

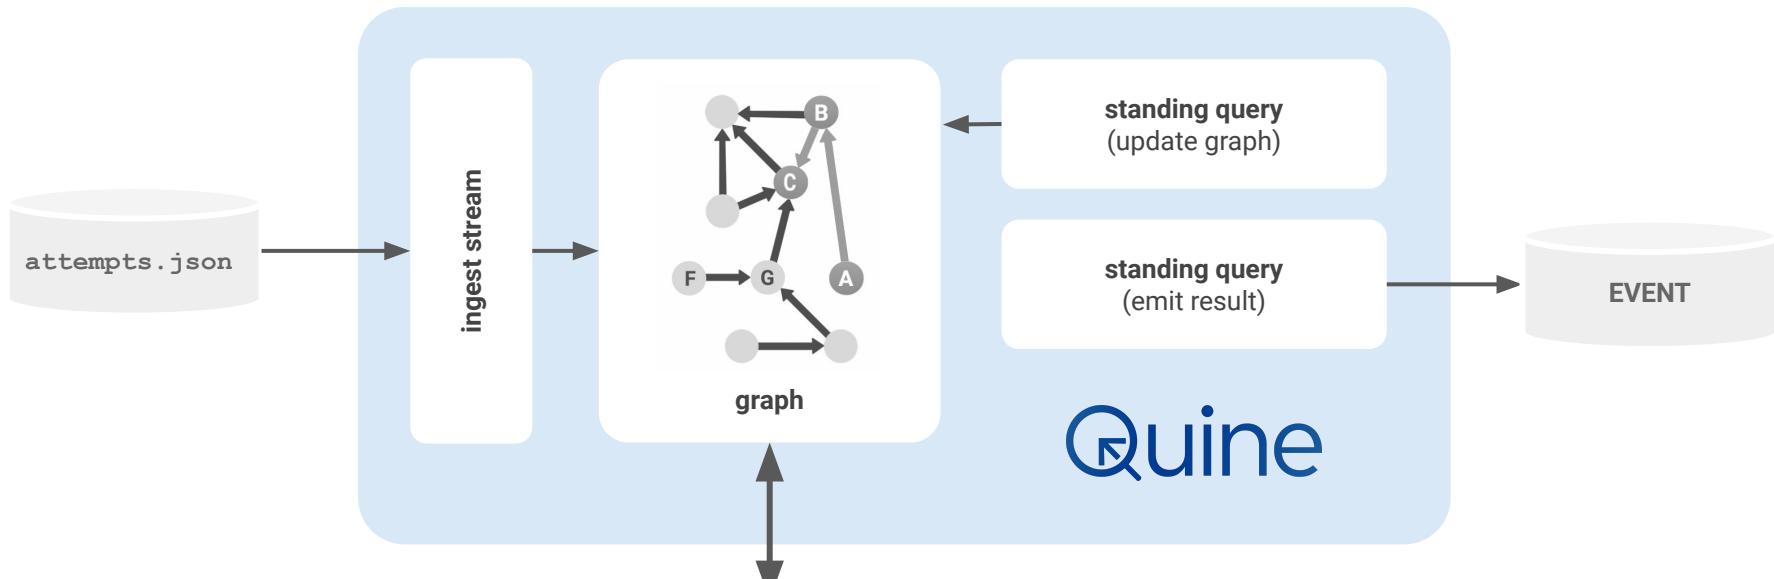
06

What's next?
Quiz, Homework, Next week



Agenda



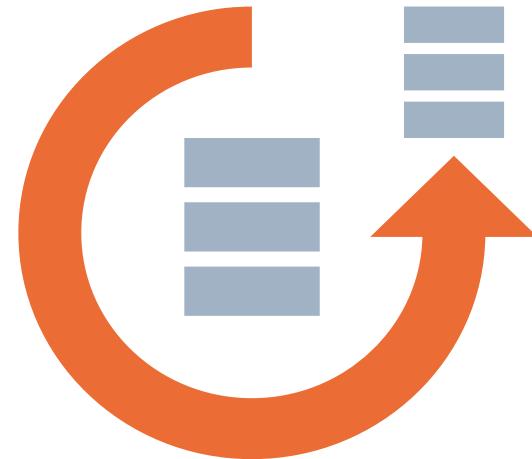


Cassandra

Architecture

Persistence Layer Options:

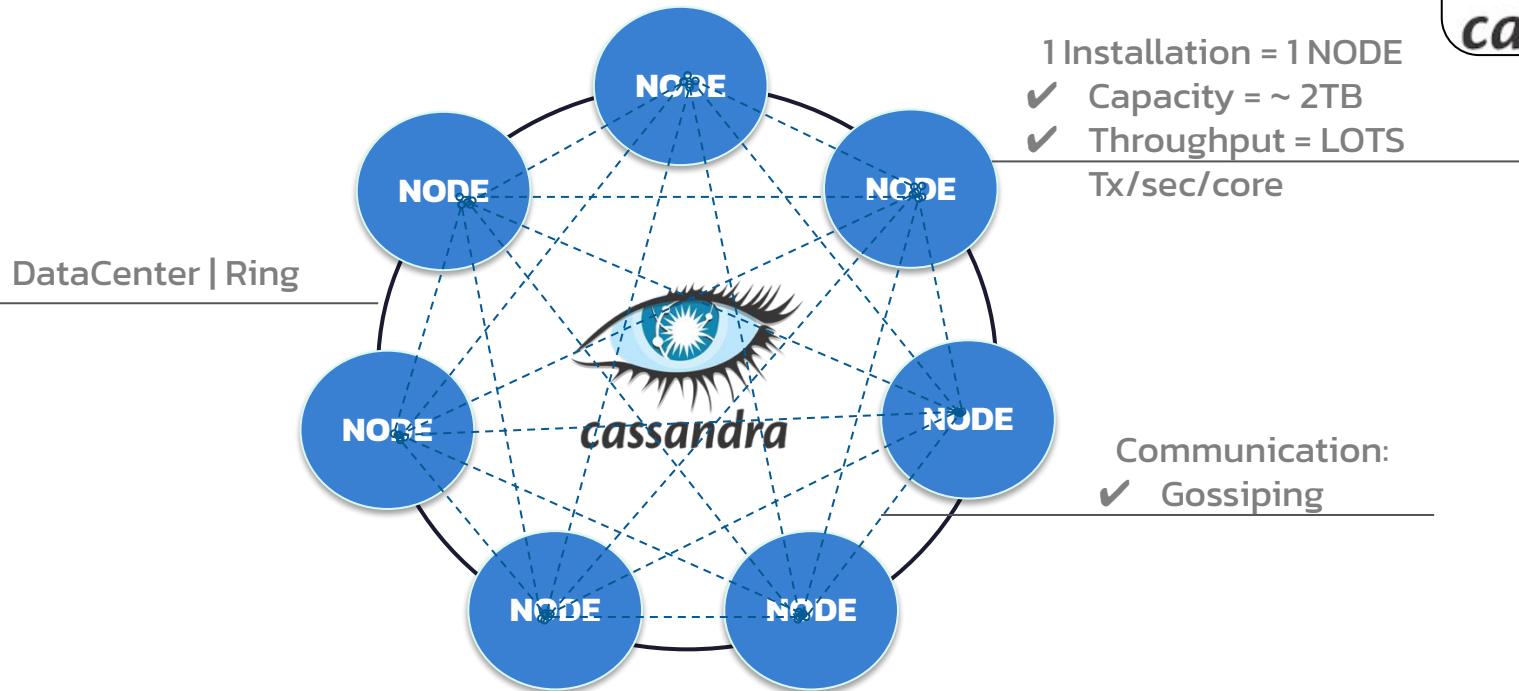
- RocksDB
- MapDB
- Apache Cassandra®



Quine storage options



Apache Cassandra®



Database: Cassandra

Astra DB:

- Cassandra as a Service!
- No operational overhead
- Geographic Awareness
- Not cloud vendor bound
- “No touch” Scalability

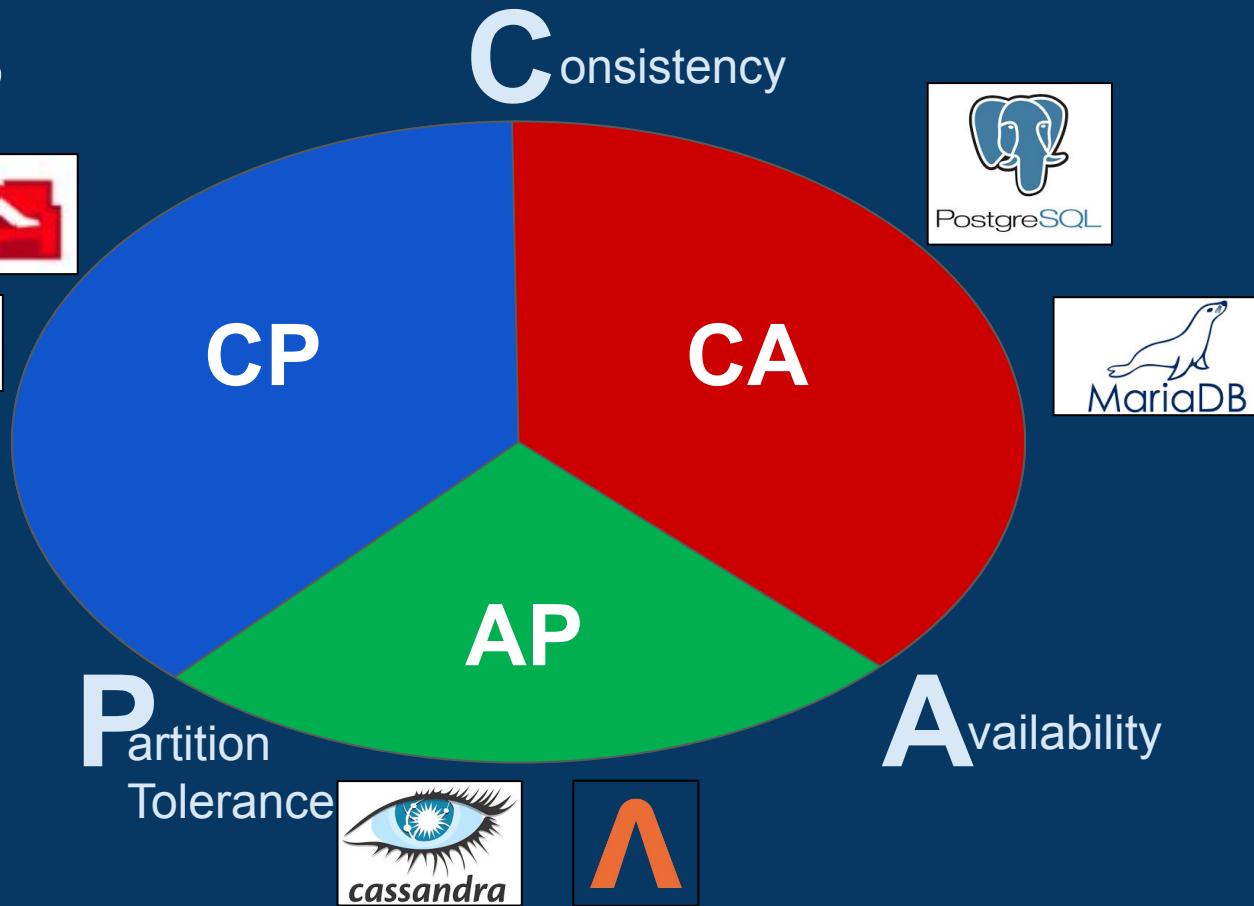


Database: Why Astra DB?

Brewer's CAP Theorem



All distributed systems try to achieve consistency, availability and partition tolerance. But realistically they can only pick two.



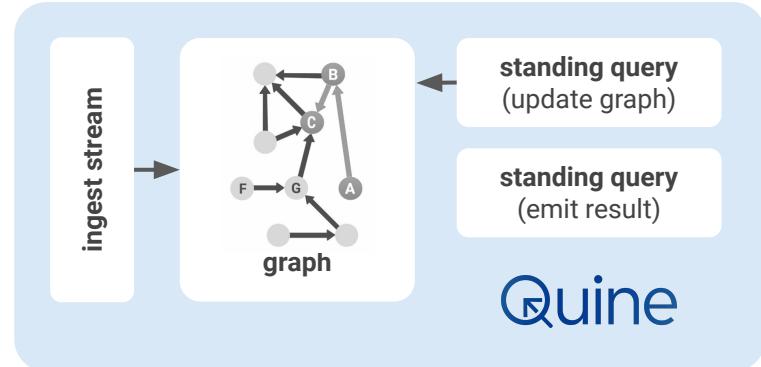
Harvest, Yield, and Scalable Tolerant Systems (1999)

CAP Theorem and Astra DB



Event Streams and Indexes

- Quine shapes event streams – data that is **continuously** generated, often in high **volumes** and at high **velocity** – into a graph.
- Traditional node indexing methods are **slow**.
- That is why we use a custom Cypher function `idFrom` which implements a consistent hash to **deterministically** locate a node ID
- Eliminates querying for node IDs
- Cornerstone concept for Quine Cypher queries



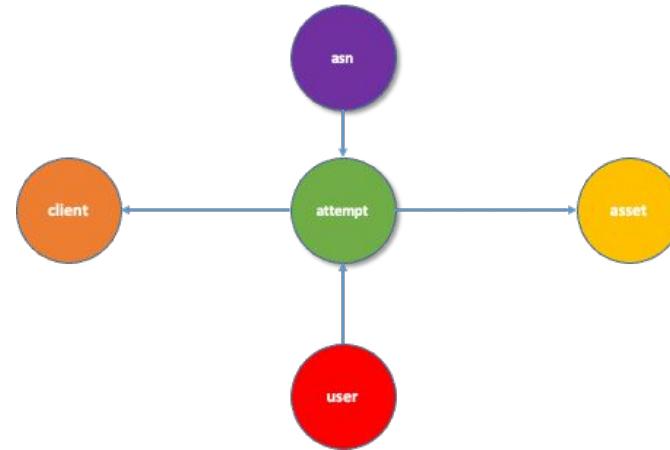
```
MATCH (n) WHERE id(n) = idFrom($that) SET n.line = $that
```

Cypher

Key Concept - `idFrom()`

Attempt Event

```
1 {
2   "schemaVersion": 1,
3   "eventId": "c7bfd702-88e3-47a5-93c5-8b699e803674",
4   "timestamp": "2022-01-01T00:08:11Z",
5   "level": "info",
6   "eventCode": 1,
7   "eventType": "user.session.auth",
8   "displayMessage": "User login",
9   "transaction": {
10     "type": "LDAP",
11     "entityId": "fd5296ee-346a-41be-b2a8-10f10319cd45",
12     "id": "760f08c0-d458-4a31-97ff-5856a1b289e6"
13   },
14   "eventTypeDetail": "user.authentication.login.success",
15   "outcome": {
16     "result": "SUCCESS",
17     "details": ""
18   },
19   "client": {
20     "zone": "VPN",
21     "asn": 12414,
22     "ipAddress": "220.209.114.209",
23     "requestUri": "/rhoncus.js?leo=faucibus&maecenas=orci&pulvinar=luctus&lobortis=et&est=ultrices&phasellus=posuere&sit=cubilia&amet=curae&erat=mauris&nulla=vierra&tempus=diam&vivamus=vitae&in=quam&felis=suspendisse&eu=potentii&apien=nullam&cursus=porttitor&vestibulum=lacus&proin=at&eu=turpis&mi=donec&nulla=posuere&ac=metus&enim=vitae&in=ipsum&tempor=aliquam",
24     "device": "unknown",
25     "userAgent": "Mozilla/5.0 (Windows NT 5.1) AppleWebKit/535.1 (KHTML, like Gecko) Chrome/14.0.792.0 Safari/535.1"
26   },
27   "user": {
28     "id": "9c3547ae-e363-4fe6-893f-87caae67d186",
29     "type": "Admin",
30     "alternateId": "Caro.Howell@acme.com",
31     "displayName": "Caro Howell"
32   }
33 }
```

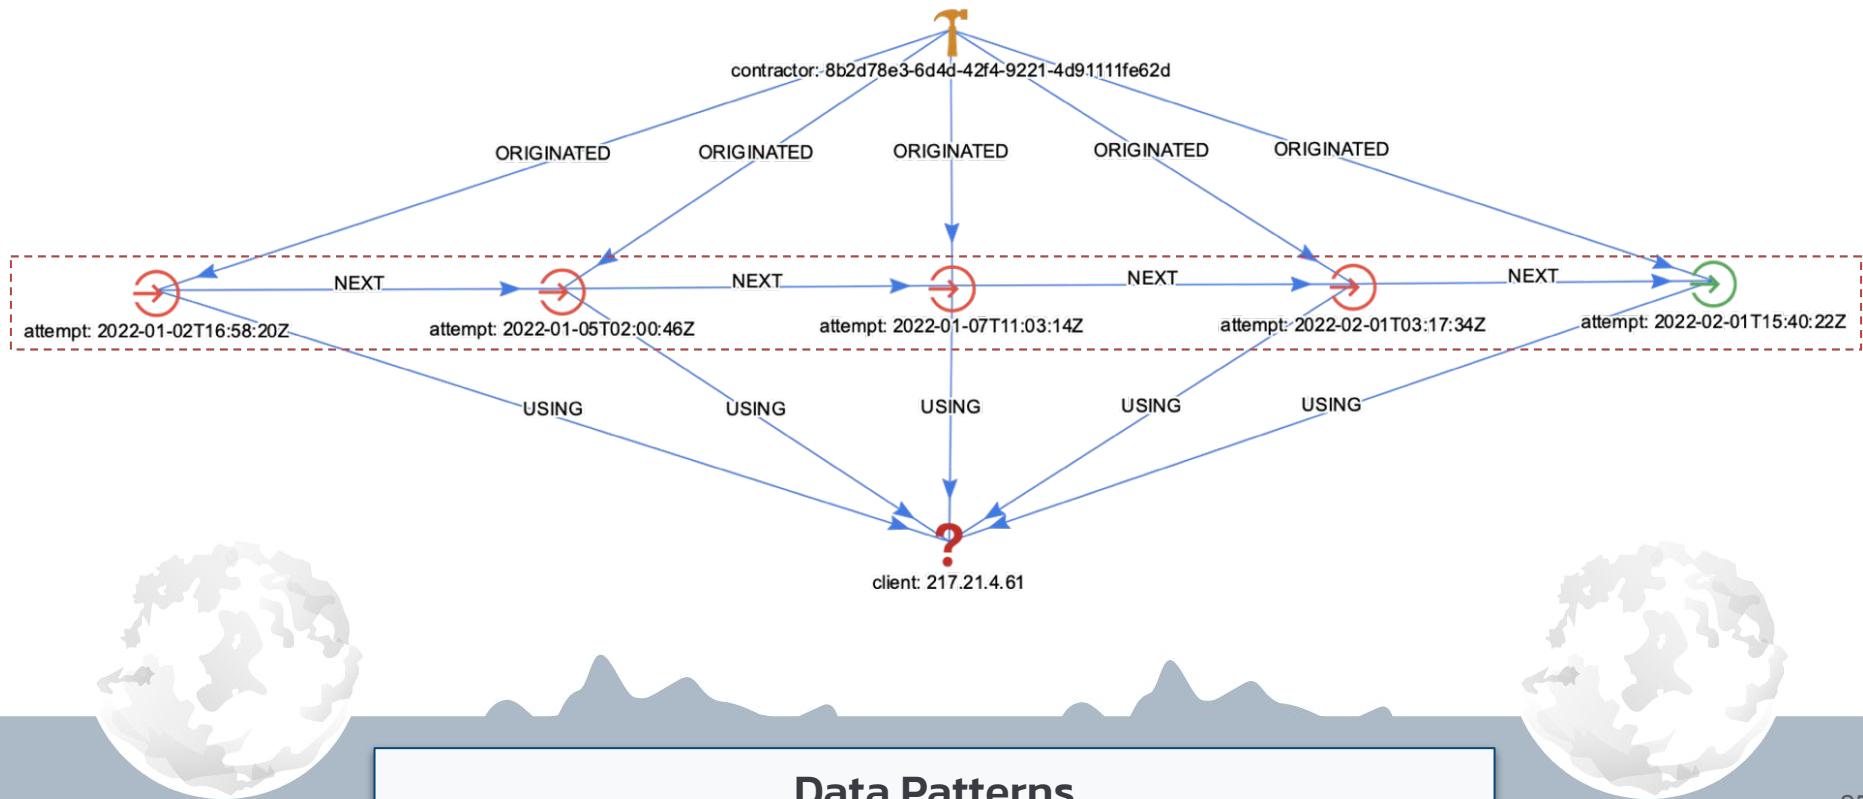


The **attempt** event is the most atomic unit. It represents a single log entry.

All nodes are materialized from the **attempt** data during ingest using Cypher to create nodes, properties, and relationships.

Event -> Ingest Stream

Four sequential **failed** attempts from the same user/client pair



```

- pattern:
  type: Cypher
  query: |-
    ///////////////////////////////////////////////////////////////////
    // Subquery to match 4 consecutive failed attempts followed by a successful attempt
    ///////////////////////////////////////////////////////////////////
    MATCH (attempt1 {outcomeResult: "FAILURE"})-[:NEXT]-(attempt2 {outcomeResult: "FAILURE"})-[:NEXT]-(attempt3 {outcomeResult: "FAILURE"})-[:NEXT]-(attempt4 {outcomeResult: "FAILURE"})-[:NEXT]-(attempt5 {outcomeResult: "SUCCESS"})
    RETURN DISTINCT id(attempt1) AS attempt1
  mode: DistinctId
  outputs:
  alert:
    type: CypherQuery
    query: |-
      MATCH (attempt1 {outcomeResult:"FAILURE"})-[:NEXT]-(attempt2 {outcomeResult:"FAILURE"})-[:NEXT]-(attempt3 {outcomeResult:"FAILURE"})-[:NEXT]-(attempt4 {outcomeResult:"FAILURE"})-[:NEXT]-(attempt5 {outcomeResult:"SUCCESS"})
      WHERE id(attempt1)=$that.data.attempt1
      RETURN 'Password Spraying Attack: ' + 'http://localhost:8080/#MATCH%20(user)-[:ORIGINATED]-%3E(attempt1%20%7BoutcomeResult:%22FAILURE%22%7D)-[:NEXT]-%3E(attempt2%20%7BoutcomeResult:%22FAILURE%22%7D)-[:NEXT]-%3E(attempt3%20%7BoutcomeResult:%22FAILURE%22%7D)-[:NEXT]-%3E(attempt4%20%7BoutcomeResult:%22FAILURE%22%7D)-[:NEXT]-%3E(attempt5%20%7BoutcomeResult:%22SUCCESS%22%7D)%20WHERE%20id(attempt1)=%22' + toString(strId(attempt1)) + '%22%20RETURN%20DISTINCT%20user%2Cattempt1%2Cattempt2%2Cattempt3%2Cattempt4%2Cattempt5' AS QuineUILink
  andThen:
    type: PrintToStandardOut

```

The Bad Pattern

01



Introduction

02



Quine streaming
graph

03



Use Cases

04



cassandra

Architecture

05



Hands on

06

What's next?
Quiz, Homework, Next week



Agenda





Jump to README steps #4 and #5

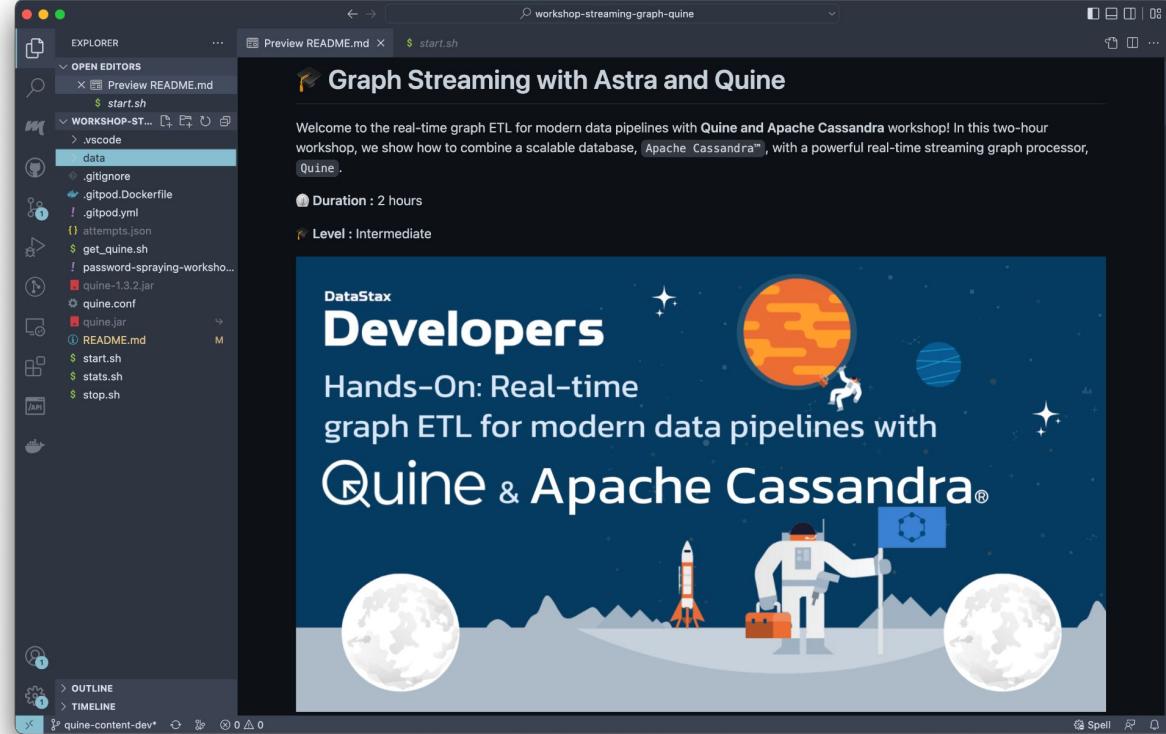
<https://github.com/datastaxdevs/workshop-streaming-graph-quine#>

[create-your-astra-db-instance](#)

<https://github.com/datastaxdevs/workshop-streaming-graph-quine#setup-quine>



Hands On



GitPod

01



Introduction

02



Quine streaming
graph

03



Use Cases

04



Architecture

05



Hands on

06

What's next?
Quiz, Homework, Next week



Agenda





Go to www.menti.com and use the code 3491 9972

Leaderboard

4821 p	spanda
4820 p	Agent X9
4775 p	fastest
4775 p	Sam
4711 p	CCedrickThePresenter
4468 p	shubham
4371 p	aaa
3895 p	vignesh
3877 p	adry
3861 p	Millie
3812 p	Puggie

2.11.07 / 2.26.05

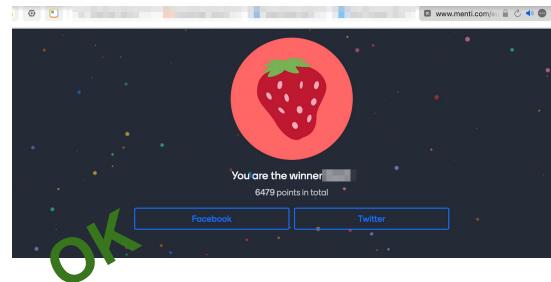
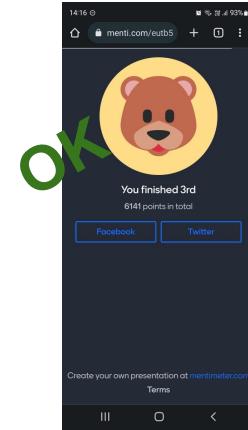


Play with us with Menti.com (new TAB)

SWAG WINNERS



Congratulations to 1st, 2nd and 3rd place
on the Menti quiz!



To claim your prize:

Take a screenshot of your Menti screen

Fill the form at dtsx.io/workshop-swag

NO!



Swag Winners!

Homework



[Complete hands-on]

1. Modify an existing recipe to do something new, and take SCREENSHOT(s).
2. Send the SCREENSHOT(s) and recipe YAML.
3. Email to Aaron and Michael.

aaron.ploetz@datastax.com,
michael@thatdot.com

Subject: Quine Homework



Cassandra Days ... in-person events are back!

- | | | | |
|---------------|---|-------------|--------|
| ● Berlin | Sept 20 | ● Hanoi | Nov 8 |
| ● London | Oct 11 | ● Jakarta | Nov 10 |
| ● Amsterdam | Oct 13 | ● Singapore | Nov 15 |
| ● Santa Clara | Nov 10 | | |
| ● Seattle | Nov 10 (thatDot/Quine speaker) | | |
| ● Houston | Nov 10 (thatDot/Quine speaker) | | |

Find out more and register at <https://www.datastax.com/events>





KubeCon

CloudNativeCon

North America 2022

Register

Attend ▾

Venue + Travel

Sponsor

Program ▾

Contact Us

View All
Events



KubeCon

CloudNativeCon

North America 2022

OCTOBER 24 – 28

DETROIT, MICHIGAN

REGISTER

VIEW THE SCHEDULE



CASSANDRA SUMMIT

MARCH 13-14, 2023 • SAN JOSE, CA

SAVE THE DATE

MCENERY CONVENTION CENTER
SAN JOSE, CALIFORNIA + VIRTUAL

- **Training day March 12**
- **In-person**
- **Virtual**
- **CFP and Reg coming soon**

Meet with us. Win a Tesla.

Let's talk DBaaS built on
Apache Cassandra®.

BOOK A MEETING NOW



*3-year Tesla Model 3 lease. NO PURCH NEC. Ends 12/31/22.

DataStax

intel.
XEON®

bit.ly/datastax_tesla

[*Terms & conditions apply](#)

DataStax

Thank You!

